

EES 3310/5310 Lab #2 Assignment, Part 2

Exercises in Data Manipulation

put your name here

Lab: Mon. Jan. 13. Due: Fri. Jan. 24

Contents

Instructions	1
Worked Example	1
Downloading CO ₂ Data from Mauna Loa Observatory	1
Exercises	9
Exercises with CO ₂ Data from the Mauna Loa Observatory	9
Exercises with Global Temperature Data from NASA	11

Instructions

This exercise starts with a worked example that shows how to download measurements of carbon dioxide from the laboratory on Mauna Loa, Hawaii, that was started by C. David Keeling in 1958, analyze the data, and make plots and tables to show the results of your analysis.

After studying the worked example, you will do further analysis and plotting using both the CO₂ data from Mauna Loa and also global temperature measurements that you will download from NASA's Goddard Institute for Space Studies.

The section "Exercises" has instructions about what to do and places where you will fill in R code, following the instructions, in order to perform these analysis.

Worked Example

Downloading CO₂ Data from Mauna Loa Observatory

In 1957, Charles David Keeling established a permanent observatory on Mauna Loa, Hawaii to make continuous measurements of atmospheric carbon dioxide. The observatory has been running ever since, and has the longest record of direct measurements of atmospheric carbon dioxide levels.

The location was chosen because the winds blow over thousands of miles of open ocean before reaching Mauna Loa, and this means the CO₂ measurements are very pure and uncontaminated by any local sources of pollution.

We can download the data from https://scrippsco2.ucsd.edu/assets/data/atmospheric/stations/in_situ_co2/monthly/monthly_in_situ_co2_mlo.csv. We can download the file and save it to the local computer using the R function `download.file`

```
download.file(mlo_url, '_data/mlo_data.csv')
```

Try opening the data file in Excel or a text editor.

The first 54 lines of the data file are comments describing the data. These comments describe the contents of each column of data and explain that this data file uses the special value -99.99 to indicate a missing value. The Mauna Loa Observatory started recording data in March 1958, so the monthly averages for January and February are missing. Other months are missing for some months in the record when the instruments that monitor CO₂ concentrations were not working properly.

The `read_csv` function from the `tidyverse` package can read the data into R and convert it into a `tibble` data structure (like a fancy data table). However, when R reads in `.csv` files, it expects column names to be on a single row, and lines 55–57 of the data file are column headings that are split across multiple rows, so R will get confused if we tell it to use those rows as column names.

Thus, we will tell `read_csv` to read this data file, but skip the first 57 lines. We will also tell it not to look for column names in the data file, so we will supply the column names, and we will tell it that whenever it sees -99.99, it should interpret that as indicating a missing value, rather than a measurement.

Finally, R can guess what kinds of data each column contains, but for this file, things work a bit more smoothly if we provide this information explicitly.

`read_csv` lets us specify the data type for each column by providing a string with one letter for each column. The letters are `i` for integer numbers, `d` for real numbers (i.e., numbers with a decimal point and fractional parts), `n` for an unspecified number, `c` for character (text) data, `l` for logical (TRUE or FALSE), `D` for calendar dates, `t` for time of day, and `T` for combined date and time.

```
mlo_data = read_csv('_data/mlo_data.csv',  
  skip = 57, # skip the first 57 rows  
  col_names = c('year', 'month', 'date.excel', 'date',  
    'co2.raw', 'co2.raw.seas',  
    'co2.fit', 'co2.fit.seas',  
    'co2.filled', 'co2.filled.seas'),  
  col_types = 'iiiddddddd',  
  # ^^^ the first three columns are integers and the next  
  # 7 are real numbers  
  na = '-99.99',  
  # ^^^ interpret -99.99 as a missing value  
)
```

Let's look at the first few rows of the data:

Here is how it looks in R:

```
head(mlo_data)

## # A tibble: 6 x 10
##   year month date.excel  date co2.raw co2.raw.seas co2.fit co2.fit.seas
##   <int> <int>   <int> <dbl>   <dbl>      <dbl>   <dbl>      <dbl>
## 1  1958     1     21200 1958.    NA        NA      NA        NA
## 2  1958     2     21231 1958.    NA        NA      NA        NA
## 3  1958     3     21259 1958.   316.     314.    316.     315.
## 4  1958     4     21290 1958.   317.     315.    317.     315.
## 5  1958     5     21320 1958.   318.     315.    318.     315.
## 6  1958     6     21351 1958.    NA        NA     317.     315.
## # ... with 2 more variables: co2.filled <dbl>, co2.filled.seas <dbl>
```

There are six different columns for the CO₂ measurements:

- `co2.raw` is the raw measurement from the instrument. The measurements began in March 1958, so there are NA values for January and February. In addition, there are missing values for some months when the instrument was not working well.
- `co2.fit` is a smoothed version of the data, which we will not use in this lab.
- `co2.filled` is the same as `co2.raw`, except that where there are missing values in the middle of the data, they have been filled in with interpolated estimates based on measurements before and after the gap.

For each of these three data series, there is also a *seasonally adjusted* version, which attempts to remove the effects of seasonal variation in order to make it easier to observe the trends.

For this lab, we will focus on the `co2.filled` data series. To keep things simple, we can use the `select` function from `tidyverse` to keep only certain columns in the tibble and get rid of the ones we don't want.

```
mlo_simple = mlo_data %>% select(year, month, date, co2 = co2.filled)

head(mlo_simple)
```

```
## # A tibble: 6 x 4
##   year month  date    co2
##   <int> <int> <dbl> <dbl>
## 1  1958     1 1958.    NA
## 2  1958     2 1958.    NA
## 3  1958     3 1958.  316.
## 4  1958     4 1958.  317.
## 5  1958     5 1958.  318.
## 6  1958     6 1958.  317.
```

Note how we renamed the `co2.filled` column to just plain `co2` in the `select` function.

We can also use the `kable()` function to format the data nicely as a table in an RMarkdown document:

```
head(mlo_simple) %>%
  knitr::kable(col.names = c(year = "Year", month = "Month", date = "Date",
                             co2 = "CO~2~ (ppm)"),
               caption = "A table of monthly CO~2~ measurements from Mauna Loa.")
```

Table 1: A table of monthly CO₂ measurements from Mauna Loa.

Year	Month	Date	CO ₂ (ppm)
1958	1	1958.041	NA
1958	2	1958.126	NA
1958	3	1958.203	315.70
1958	4	1958.288	317.45
1958	5	1958.370	317.51
1958	6	1958.455	317.24

Now, let's plot this data:

```
ggplot(mlo_simple, aes(x = date, y = co2)) +
  # ~~~ The ggplot command specifies which data to plot and the aesthetics that
  # define which variables to use for the x and y axes.
  geom_line() +
  # ~~~ The geom_line() command says to plot lines connecting the data points
  labs(x = "Year", y = "CO2 concentration (ppm)",
       title = "Measured CO2 from Mauna Loa Observatory")
  # ~~~ The labs() command tells ggplot what names to use in labeling the axes
  # and the title for the plot.

# Earlier in this .Rmd file, I called set_theme(theme_bw(base_size = 15))
# to set the default plot style. If you call ggplot() without this,
# you will get a different style, but you can either call theme_set
# or you can add a theme specification (such as
# "+ theme_bw(base_size = 15)")
# to the end of the sequence of plotting commands in order to
# apply a specific style to an individual plot.
```

I created a caption for the figure caption by adding the following specification to the header of the R code chunk in the RMarkdown document:

```
fig.cap="Monthly CO~2~ measurements from Mauna Loa."
```

Notice the seasonal variation in CO₂. Every year, there is a large cycle of CO₂, but underneath is a




Figure 1: Monthly CO₂ measurements from Mauna Loa.

gradual and steady increase from year to year. If we wanted to look at the trend without the seasonal variation, we could use the `co2.filled.seas` column of the original tibble, but instead, let's look at how we might estimate this ourselves.

The seasonal cycle is 12 months long and it repeats every year. This means that if we average the values in our table over a whole year, this cycle should average out. We can do this by creating a new column `annual` where every row represents the average over a year centered at that row (technically, all the months from 5 months before through six months after that date).

To do this, we use the function `rollapply` from the `zoo` package, as shown below. The `rollapply` function allows you to take a series of data (such as monthly CO₂ measurements) and at each point, apply a function to the data within a “window” that includes a certain number of points before and after the point in question.

Here, we apply the mean function (`FUN = mean`) to take the average, and we define the “window” to be the 12 points centered on the point in question. This means that for each month in our data series, `rollapply` takes the average of the 12 measurements centered on that month (technically, the month, the five months before, and the six months after). You could also specify `align = "left"` to take the 12 months starting with the given month, or `align = "right"` to take the 12 months ending with the given month.

We also need to tell `rollapply` what to do with the months at the beginning of the series that don't have five months of data before them and points at the end of the series that don't have six months after them. `fill = NA` means to set those points to NA, which is a special value R uses to indicate missing values (NA means “not available”).

```
mlo_simple %>% mutate(annual = rollapply(data = co2, width = 12, FUN = mean,
                                         fill = NA, align = "center")) %>%

  ggplot(aes(x = date)) +
  geom_line(aes(y = co2), color = "darkblue", size = 0.1) +
  geom_line(aes(y = annual), color = "black", size = 0.5) +
  labs(x = "Year", y = "CO2 concentration (ppm)",
       title = "Measured and Seasonally Adjusted CO2")
```

But wait: we might want a legend to tell the reader what each colored line represents. We can create new aesthetics for the graph mapping to do this:

```
mlo_simple %>% mutate(annual = rollapply(data = co2, width = 12, FUN = mean,
                                         fill = NA, align = "center")) %>%

  ggplot(aes(x = date)) +
  geom_line(aes(y = co2, color = "Raw"), size = 0.1) +
  geom_line(aes(y = annual, color = "12-month average"), size = 0.5) +
  scale_color_manual(values = c("Raw" = "darkblue",
                               "12-month average" = "black"),
                    name = "Smoothing") +
  labs(x = "Year", y = "CO2 concentration (ppm)",
       title = "Measured and Seasonally Adjusted CO2")
```

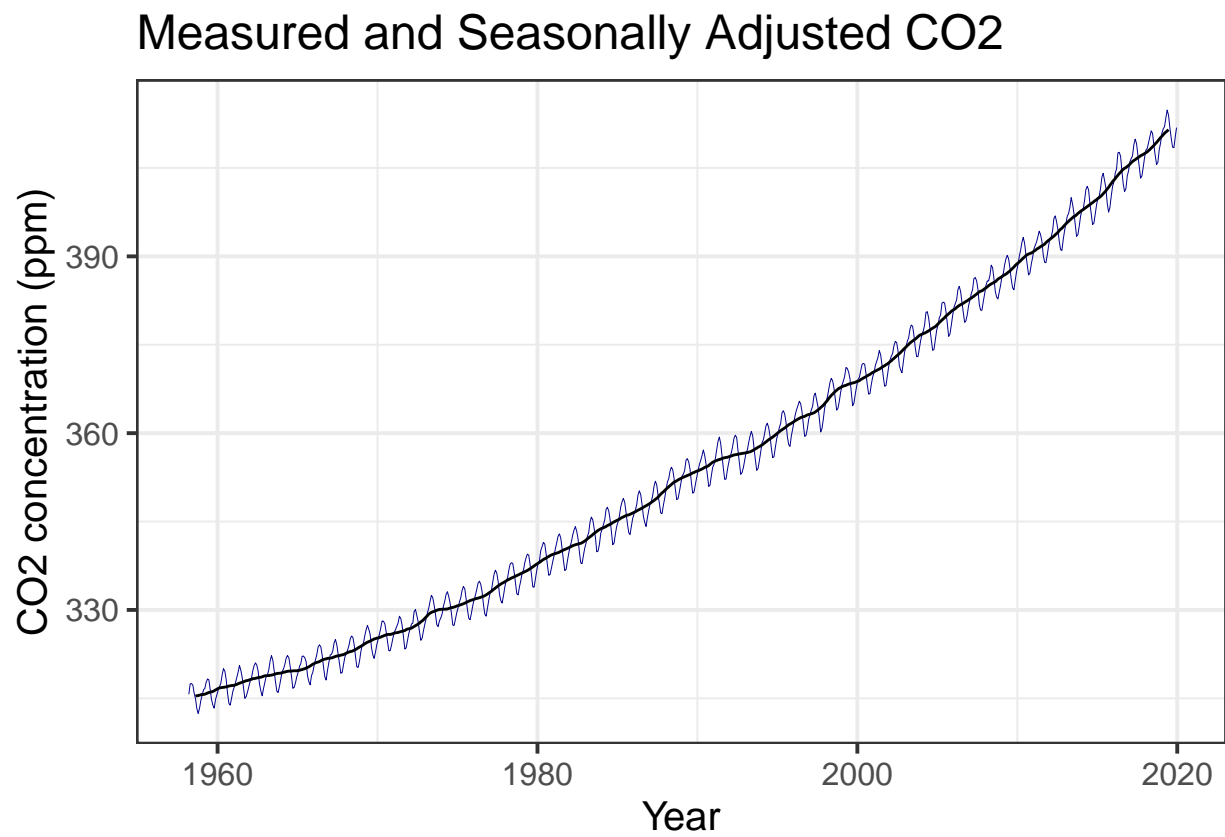


Figure 2: Raw and seasonally adjusted measurements of atmospheric CO₂, from Mauna Loa.




Figure 3: Raw and seasonally adjusted measurements of atmospheric CO₂, from Mauna Loa, with a legend identifying the different lines.

We can also analyze this data to estimate the average trend in CO₂. We use the `lm` function in R to fit a straight line to the data, and we use the `tidy` function from the `broom` package to print the results of the fit nicely.

R has many powerful functions to fit data, but here we will just use a very simple one. We specify the linear relationship to fit using R's formula language. If we want to tell R that we think `co2` is related to `date` by the linear relationship $co2 = a + b \times date$, then we write the formula `co2 ~ date`. The intercept is implicit, so we don't have to spell it out.

```
trend = lm(co2 ~ date, data = mlo_simple)

library(broom)

tidy(trend)

## # A tibble: 2 x 5
##   term          estimate std.error statistic p.value
##   <chr>          <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept) -2763.      17.7      -156.      0
## 2 date           1.57     0.00889     176.      0
```

This shows us that the trend is for CO₂ to rise by 1.57 ppm per year, with an uncertainty of plus or minus 0.009.

We can also plot a linear trend together with the data:

```
mlo_simple %>%
  ggplot(aes(x = date, y = co2)) +
  geom_line() +
  geom_smooth(method = 'lm') +
  labs(x = "Year", y = "CO2 concentration (ppm)",
       title = "Measured CO2 and Linear Fit")
```

Exercises

Exercises with CO₂ Data from the Mauna Loa Observatory

Using the `select` function, make a new data tibble called `mlo_seas`, from the original `mlo_data`, which only has two columns: `date` and `co2.seas`, where `co2.seas` is a renamed version of `co2.filled.seas` from the original tibble.

```
# put your R code here
```

Now plot this with `co2.seas` on the y axis and `date` on the x axis, and a linear fit:

```
# put your R code here
# remember to use geom_smooth to include a linear fit.
```




Figure 4: Trend in atmospheric CO₂.

Now fit a linear function to find the annual trend of `co2.seas`. Save the results of your fit in a variable called `trend.seas`.

```
# put your R code here to set trend.seas
```

Compare the trend you fit to the raw `co2.filled` data to the trend you fit to the seasonally adjusted data.

Exercises with Global Temperature Data from NASA

We can also download a data set from NASA's Goddard Institute for Space Studies (GISS), which contains the average global temperature from 1880 through the present.

The URL for the data file is https://data.giss.nasa.gov/gistemp/tabledata_v3/GLB.Ts+dSST.csv

Download this file and save it in the directory `_data/global_temp_land_sea.csv`.

```
# Put your R code here
```

- Open the file in Excel or a text editor and look at it.
- Unlike the CO₂ data file, this one has a single line with the data column names, so you can specify `col_names=TRUE` in `read_csv` instead of having to write the column names manually.
- How many lines do you have to tell `read_csv` to skip?
- `read_csv` can automatically figure out the data types for each column, so you don't have to specify `col_types` when you call `read_csv`
- This file uses `***` to indicate missing values instead of `-99.99`, so you will need to specify `na="***"` in `read_csv`.

For future reference, if you have a file that uses multiple different values to indicate missing values, you can give a vector of values to `na` in `read_csv`: `na = c('***', '-99.99', 'NA', ' ')` would tell `read_csv` that if it finds any of the values `"***"`, `"-99.99"`, `"NA"`, or just a blank with nothing in it, any of those would correspond to a missing value, and should be indicated by `NA` in R.

Now read the file into R, using the `read_csv` function, and assign the resulting tibble to a variable `giss_temp`

```
# Put your R code here to call read_csv and read "global_temp_land_sea.csv"
```

```
# Then, show the first 5 lines of giss_temp:  
# head(giss_temp, 5)
```

Something is funny here: Each row corresponds to a year, but there are columns for each month, and some extra columns called “J-D”, “D-N”, “DJF”, “MAM”, “JJA”, and “SON”. These stand for average values for the year from January through December, the year from the previous December through November, and the seasonal averages for Winter (December, January, and February), Spring

(March, April, and May), Summer (June, July, and August), and Fall (September, October, and November).

The temperatures are recorded not as the thermometer reading, but as *anomalies*. If we want to compare how temperatures are changing in different seasons and at different parts of the world, raw temperature measurements are hard to work with because summer is hotter than winter and Texas is hotter than Alaska, so it becomes difficult to compare temperatures in August to temperatures in January, or temperatures in Texas to temperatures in Alaska and tell whether there was warming.

To make it easier and more reliable to compare temperatures at different times and places, we define anomalies: The temperature anomaly is the difference between the temperature recorded at a certain location during a certain month and a baseline reference value, which is the average temperature for that month and location over a period that is typically 30 years.

The GISS temperature data uses a baseline reference period of 1951–1980, so for instance, the temperature anomaly for Nashville in July 2017 would be the monthly average temperature measured in Nashville during July 2017 minus the average of all July temperatures measured in Nashville from 1951–1980.

The GISS temperature data file then averages the temperature anomalies over all the temperature-measuring stations around the world and reports a global average anomaly for every month from January 1880 through the latest measurements available (currently, July 2017).

Let's focus on the months only. Use `select` to select just the columns for “Year” and January through December (if you are selecting a consecutive range of columns between “Foo” and “Bar”, you can call `select(Foo:Bar)`). Save the result in a variable called `giss_monthly`

```
# put your R code here
```

Next, it will be difficult to plot all of the data if the months are organized as columns. What we want is to transform the data tibble into one with three columns: “year”, “month”, and “anomaly”. We can do this easily using the `gather` function from the tidyverse package: `gather(df, key = month, value = anomaly, -Year)` or `df %>% gather(key = month, value = anomaly, -Year)` will gather all of the columns except Year (the minus sign in `select` or `gather` means to include all columns except the ones indicated with a minus sign) and:

- Make a new tibble with three columns: “Year”, “month”, and “anomaly”
- For each row in the original tibble, make rows in the new tibble for each of the columns “Jan” through “Dec”, putting the name of the column in “month” and the anomaly in “anomaly”.

Here is an example of using `gather`, using the built-in data set `presidents`, which lists the quarterly approval ratings for U.S. presidents from 1945–1974:

```
df = matrix(presidents, ncol=4, byrow = TRUE) %>%  
  as_tibble() %>% set_names(paste0("Q", 1:4)) %>% mutate(year = 1944 + seq(n()))
```

```
## Warning: `as_tibble.matrix()` requires a matrix with column names or a `.name_repair`  
## This warning is displayed once per session.
```

```
print("First 10 rows of df are")
```

```
## [1] "First 10 rows of df are"
```

```
print(head(df, 10))
```

```
## # A tibble: 10 x 5
##       Q1     Q2     Q3     Q4  year
##   <dbl> <dbl> <dbl> <dbl> <dbl>
## 1    NA    87    82    75  1945
## 2    63    50    43    32  1946
## 3    35    60    54    55  1947
## 4    36    39    NA    NA  1948
## 5    69    57    57    51  1949
## 6    45    37    46    39  1950
## 7    36    24    32    23  1951
## 8    25    32    NA    32  1952
## 9    59    74    75    60  1953
## 10   71    61    71    57  1954
```

For each year, the table has a column for the year and four columns (Q1 ... Q4) that hold the quarterly approval ratings for the president in that quarter. Now we want to gather these data into three columns: one column for the year, one column to indicate the quarter, and one column to indicate the approval rating.

We do this with the `gather` function from the `tidyverse` package.

```
““{ r gather_example} dfg = df %>% gather(key = quarter, value = approval, -year) %>% ar-
range(year, quarter) # the gather command gathers the data into tidy columns: # “key = quarter” tells
gather to create a column called “quarter” and store # the names of the columns that are gathered
there. # “value = approval” tells gather to create a column called “approval” and # store the values
from the columns there. # “-year” tells gather NOT to gather the column “year”. # So the approval
ratings from the second quarter of 1960 will be stored in # a row where the column year contains
1960, quarter contains “Q2”, and # approval contains the approval rating. # # the arrange command
sorts the rows of the resulting tibble to put # the years in ascending order, from 1945 to 1971, and
within each year, # sort the quarters in alphabetical order from Q1 to Q4
```

```
head(dfg) # print the first few rows of the tibble.
```

Now you try to do the same thing to:

- * First select just the columns of ``giss_monthly`` for the year and the individual months.
- * Next, gather all the months together, so there will be three columns: one for the year, one for the name of the month, and one for the temperature anomaly in that month.

* Store the result in a new variable called ``giss_g``

```
```r
```

```
put your R code here
```

Remember how the CO<sub>2</sub> data had a column `date` that had a year plus a fraction that corresponded to the month, so June 1960 was 1960.4548?

Here is a trick that lets us do the same for the `giss_g` data set. R has a data type called `factor` that it uses for managing categorical data, such as male versus female, Democrat versus Republican, and so on. Categorical factors have a textual label, but are silently represented as integer numbers. Normal factors don't have a special order, so R sorts the values alphabetically. However, there is another kind of factor called an ordered factor, which allows us to specify the order of the values.

We can use a built-in R variable called `month.abb`, which is a vector of abbreviations for months.

The following command will convert the `month` column in `giss_g` into an ordered factor that uses the integer values 1, 2, ..., 12 to stand for "Jan", "Feb", ..., "Dec", and then uses those integer values to create a new column, `date` that holds the fractional year, just as the `date` column in `mlo_data` did:

```
giss_g = giss_g %>%
 mutate(month = ordered(month, levels = month.abb),
 date = Year + (as.integer(month) - 0.5) / 12) %>%
 arrange(date)`
```

In the code above, `ordered(month, levels = month.abb)` converts the variable `month` from a character (text) variable that contains the name of the month to an ordered factor that associates a number with each month name, such that 'Jan' = 1 and 'Dec' = 12.

Then we create a new column called `date` to get the fractional year corresponding to that month. We have to explicitly convert the ordered factor into a number using the function `as.integer()`, and we subtract 0.5 because the time that corresponds to the average temperature for the month is the middle of the month.

Below, use code similar to what I put above to add a new `date` column to `giss_g`.

```
put your R code here
```

Now plot the monthly temperature anomalies versus date:

```
put your R code here
ggplot() +
specify the data and the mappings of variables to plot aesthetics
geom_XXX() +
put the geometries (lines, points, etc) that you want to plot
labs() +
```

```
label the axes
... # put any other characteristics here.
```

That plot probably doesn't look like much, because it's very noisy. Use the function `rollapply` from the package `zoo` to create new columns in `giss_g` with 12-month and 10-year (i.e., 120-month) rolling averages of the anomalies.

Make a new plot in which you plot a thin blue line for the monthly anomaly (use `geom_line(aes(y = anomaly), color = "blue", alpha = 0.3, size = 0.1)`; `alpha` is an optional specification for transparency where 0 means invisible (completely transparent) and 1 means opaque), a medium dark green line for the one-year rolling average, and a thick dark blue line for the ten-year rolling average.

```
put your R code here
giss_g %>%
mutate(...) %>%
^^^ fill in code for "..." in "mutate()" to add a columns called
"smooth.1" for the one-year smoothed anomaly
and "smooth.10" for the ten-year smoothed anomaly.
ggplot(aes(...)) +
^^^ Then we send the result of mutate to ggplot() where it becomes the
data to plot.
Add code for the aesthetics ("...") to map variables to aesthetics.
geom_line(aes(y = anomaly), alpha = 0.3, size = 0.1) +
^^^ plot a thin blue line with the un-smoothed anomaly
geom_line(...) +
^^^ Now add a medium "darkgreen" line for the one-year smoothed data
geom_line(...) +
^^^ And a thick "darkblue" line for the ten-year smoothed data
labs(...) + # Label your axes
... # add any other plot specifications you need.
```

The graph shows that temperature didn't show a steady trend until starting around 1970, so we want to isolate the data starting in 1970 and fit a linear trend to it.

To select only rows of a tibble that match a condition, we use the function `filter` from the `tidyverse` package:

`data_subset = df %>% filter( conditions )`, where `df` is your original tibble and `conditions` stands for whatever conditions you want to apply. You can make a simple condition using equalities or inequalities:

- `data_subset = df %>% filter( month == "Jan" )` to select all rows where the month is "Jan"
- `data_subset = df %>% filter( month != "Aug" )` to select all rows where the month is not August.

- `data_subset = df %>% filter( month %in% c("Sep", "Oct", "Nov"))` to select all rows where the month is one of “Sep”, “Oct”, or “Nov”.
- `data_subset = df %>% filter(year >= 1945)` to select all rows where the year is greater than or equal to 1945.
- `data_subset = df %>% filter(year >= 1951 & year <= 1980 )` to select all rows where the year is between 1951 and 1980, inclusive.
- `data_subset = df %>% filter(year >= 1951 | month == "Mar")` to select all rows where the year is greater than or equal to 1951 or the month is “Mar”. this will give all rows from January 1951 onward, plus all rows before 1951 where the month is March.

Below, create a new variable `giss_recent` and assign it a subset of `giss_g` that has all the data from January 1970 through the present. Fit a linear trend to the monthly anomaly and report it.

What is the average change in temperature from one year to the next?

```
put your R code here
create giss_recent
fit a linear trend to the data using lm()
print the coefficients of the trend.
```

## Did Global Warming Stop after 1998?

It is a common skeptic talking point that global warming stopped in 1998. In years with strong El Niños, global temperatures tend to be higher and in years with strong La Niñas, global temperatures tend to be lower. We will discuss why later in the semester.

The year 1998 had a particularly strong El Niño, and the year set a record for global temperature that was not exceeded for several years. Indeed, compared to 1998, it might look as though global warming paused for many years.

We will examine whether this apparent pause has scientific validity.

To begin with, we will take the monthly GISS temperature data and convert it to annual average temperatures, so we can deal with discrete years, rather than separate temperatures for each month.

We do this with the `group_by` and `summarize` functions.

We also want to select only recent data, so we arbitrarily say we will look at temperatures starting in 1979, which gives us 19 years before the 1998 El Niño.

We don’t have a full year of data for 2017, so we want to discard that because we won’t get a full year average from it.

If we go back to the original `giss_g` data tibble, run the following code:

```
When you are ready to run the code below, you can un-comment it in RStudio
by highlighting the lines you want to uncomment and pressing
Ctrl + Shift + C on Windows or Cmd + Shift + C on Mac.
(You can also comment out code using the same key combination: it works
```



```
as a toggle between commented and un-commented).
#
giss_annual = giss_g %>%
filter(year >= 1979 & year < 2017) %>%
group_by(year) %>%
summarize(anomaly = mean(anomaly)) %>%
ungroup() %>%
mutate(date = year + 0.5)
```

This code groups the giss data by the year, so that one group will have January–December 1979, another will have January–December 1980, and so forth.

Then we replace the groups of 12 rows for each year (each row represents one month) with a single row that represents the average of those 12 months.

It is important to tell R to ungroup the data after we're done working with the groups.

Finally, we set date to year + 0.5 because the average of a year corresponds to the middle of the year, not the beginning.

Now, let's introduce a new column after, which indicates whether the data is after the 1998 El Niño:

Now plot the data and color the points for 1998 and afterward dark red to help us compare before and after 1998.

```
ggplot(giss_annual, aes(x = date, y)) +
geom_line(size = 1) +
geom_point(aes(color = year >= 1998), size = 1) +
scale_color_manual(values = c(TRUE = "darkred", FALSE = "darkblue"),
guide = "none") +
~~~ color "before" points dark blue, "after" points dark red.
guide = "none" tells ggplot not to show a legend explaining the colors.
labs(x = "Year", y = "Temperature Anomaly")
```

Does it look as though the red points are not rising as fast as the blue points?

Let's just plot the data from 1998 on:

```
Put your R code here
Filter the giss_annual data to select only the years >= 1998
plot the data
```

Now how does it look?

Let's use the filter function to break the data into two different tibbles: giss\_before will have the data from 1979–1998 and the other, giss\_after will have the data from 1998 onward (note that the year 1998 appears in both tibbles).

```
Put your R code here
```

Now use `lm` to fit a linear trend to the temperature data in `giss_before` (from 1979–1998) and assign it to a variable `giss_trend`.

Next, add a column `timing` to each of the split data sets and set the value of this column to “Before” for `giss_before` and “After” for `giss_after`.

```
Put your R code here
```

Now, combine the two tibbles into one tibble:

```
giss_combined = bind_rows(giss_before, giss_after)
```

Now let’s use `ggplot` to plot `giss_combined`:

- Aesthetic mapping:
  - Use the date column for the *x* variable.
  - Use the anomaly column for the *y* variable.
  - Use the `timing` column to set the color of plot elements
- Plot both lines and points.
  - Set the size of the lines to 1
  - Set the size of the points to 2
- Use the `scale_color_manual` function to set the color of “Before” to “darkblue” and “After” to “darkred”
- Use `geom_smooth(data = giss_before, method="lm", color = "blue", fill = "blue", alpha = 0.2, fullrange = TRUE)` to show a linear trend that is fit just to the `giss_before` data.

```
Put your R code here.
```

Try this with the parameter `fullrange` set to `TRUE` and `FALSE` in the `geom_smooth` function. What is the difference?

What this plot shows is the full data set, and a linear trend that is fit just to the “before” data. The trend line shows both the best fit for a trend (that’s the solid line) and the range of uncertainty in the fit (that’s the light blue shaded area around the line).

If the temperature trend changed after 1998 (e.g., if the warming paused, or if it reversed and started cooling) then we would expect the temperature measurements after 1998 to fall predominantly below the extrapolated trend line, and our confidence that the trend had changed would depend on the number of points that fall below the shaded uncertainty range.

How many of the red points fall below the trend line?

**Answer:** *put your answer here.*

How many of the red points fall above the trend line?

**Answer:** *put your answer here.*

What do you conclude about whether global warming paused or stopped after 1998?

**Answer:** *put your answer here.*